

House Price Prediction

2024-04-09

About the Data

The data was collected from the website 'Kaggle', which contains thousands of datasets used for training predictive algorithms. The specific dataset we are using contains training and testing data fortunately named "train" and "test", which contains multiple different variables such as area, year built, number of rooms, and more used to predict the price of a house.

Link to dataset on Kaggle: <https://www.kaggle.com/competitions/house-prices-advanced-regression-techniques/data>

Loading Packages

```
library(tidyverse)

## Warning: package 'tidyverse' was built under R version 4.2.3
## Warning: package 'ggplot2' was built under R version 4.2.3
## Warning: package 'tibble' was built under R version 4.2.3
## Warning: package 'tidyr' was built under R version 4.2.3
## Warning: package 'readr' was built under R version 4.2.3
## Warning: package 'purrr' was built under R version 4.2.3
## Warning: package 'dplyr' was built under R version 4.2.3
## Warning: package 'stringr' was built under R version 4.2.3
## Warning: package 'forcats' was built under R version 4.2.3
## Warning: package 'lubridate' was built under R version 4.2.3

## — Attaching core tidyverse packages ————— tidyverse
2.0.0 —
## ✓ dplyr      1.1.4      ✓ readr      2.1.5
## ✓ forcats   1.0.0      ✓ stringr    1.5.1
## ✓ ggplot2   3.5.0      ✓ tibble     3.2.1
## ✓ lubridate 1.9.3      ✓ tidyr      1.3.1
## ✓ purrr     1.0.2
## — Conflicts —————
```

```
tidyverse_conflicts() —
## ✖ dplyr::filter() masks stats::filter()
## ✖ dplyr::lag() masks stats::lag()
## ⓘ Use the conflicted package (<http://conflicted.r-lib.org/>) to force all
conflicts to become errors

library(corrplot)

## Warning: package 'corrplot' was built under R version 4.2.3
## corrplot 0.92 loaded

library(lubridate)
library(ggplot2)
library(readr)
library(caTools)

## Warning: package 'caTools' was built under R version 4.2.3

library(GGally)

## Warning: package 'GGally' was built under R version 4.2.3
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2

library(caret)

## Warning: package 'caret' was built under R version 4.2.3
## Loading required package: lattice
## Warning: package 'lattice' was built under R version 4.2.3
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##   lift

library(leaps)

## Warning: package 'leaps' was built under R version 4.2.3

library(gridExtra)

## Warning: package 'gridExtra' was built under R version 4.2.3
##
## Attaching package: 'gridExtra'
##
## The following object is masked from 'package:dplyr':
```

```
##
##      combine
```

Reading the data and understanding it

```
train_data = read.csv("train.csv")
test_dataa = read.csv("test.csv")
```

Converting all character columns to factor:

```
train_data <- as.data.frame(unclass(train_data), stringsAsFactors = TRUE)
test_dataa <- as.data.frame(unclass(test_dataa), stringsAsFactors = TRUE)
```

Now, lets view the first row of the training set:

```
head(train_data,1)

##   Id MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape
LandContour
## 1 1           60        RL           65    8450   Pave  <NA>      Reg
Lvl
##   Utilities LotConfig LandSlope Neighborhood Condition1 Condition2
BldgType
## 1   AllPub     Inside      Gtl      CollgCr      Norm      Norm
1Fam
##   HouseStyle OverallQual OverallCond YearBuilt YearRemodAdd RoofStyle
RoofMatl
## 1   2Story           7           5     2003      2003      Gable
CompShg
##   Exterior1st Exterior2nd MasVnrType MasVnrArea ExterQual ExterCond
Foundation
## 1   VinylSd      VinylSd   BrkFace      196      Gd      TA
PConc
##   BsmtQual BsmtCond BsmtExposure BsmtFinType1 BsmtFinSF1 BsmtFinType2
## 1      Gd      TA      No      GLQ      706      Unf
##   BsmtFinSF2 BsmtUnfSF TotalBsmtSF Heating HeatingQC CentralAir Electrical
## 1      0      150      856   GasA      Ex      Y      SBrkr
##   X1stFlrSF X2ndFlrSF LowQualFinSF GrLivArea BsmtFullBath BsmtHalfBath
FullBath
## 1      856      854           0      1710           1           0
2
##   HalfBath BedroomAbvGr KitchenAbvGr KitchenQual TotRmsAbvGrd Functional
## 1      1           3           1      Gd           8      Typ
##   Fireplaces FireplaceQu GarageType GarageYrBltd GarageFinish GarageCars
## 1      0      <NA>   Attchd      2003      RFn      2
##   GarageArea GarageQual GarageCond PavedDrive WoodDeckSF OpenPorchSF
## 1      548      TA      TA      Y      0      61
##   EnclosedPorch X3SsnPorch ScreenPorch PoolArea PoolQC Fence MiscFeature
## 1      0      0      0      0      <NA> <NA>      <NA>
##   MiscVal MoSold YrSold SaleType SaleCondition SalePrice
## 1      0      2    2008      WD      Normal      208500
```

From our first row of data, we can already see some columns such as “Alley” have missing values classified as NA. We also notice the column “ID” is essentially useless because it does not provide any descriptions to what the house may be like, it is essentially an identity column.

```
str(train_data)

## 'data.frame':    1460 obs. of  81 variables:
##  $ Id             : int   1 2 3 4 5 6 7 8 9 10 ...
##  $ MSSubClass     : int   60 20 60 70 60 50 20 60 50 190 ...
##  $ MSZoning       : Factor w/ 5 levels "C (all)","FV",...: 4 4 4 4 4 4 4 4 5
##  $ LotFrontage    : int   65 80 68 60 84 85 75 NA 51 50 ...
##  $ LotArea        : int  8450 9600 11250 9550 14260 14115 10084 10382 6120
##  $ Street         : Factor w/ 2 levels "Grvl","Pave": 2 2 2 2 2 2 2 2 2
##  $ Alley          : Factor w/ 2 levels "Grvl","Pave": NA NA NA NA NA NA NA
##  $ LotShape       : Factor w/ 4 levels "IR1","IR2","IR3",...: 4 4 1 1 1 1 4 1
##  $ LandContour    : Factor w/ 4 levels "Bnk","HLS","Low",...: 4 4 4 4 4 4 4 4
##  $ Utilities      : Factor w/ 2 levels "AllPub","NoSeWa": 1 1 1 1 1 1 1 1
##  $ LotConfig      : Factor w/ 5 levels "Corner","CulDSac",...: 5 3 5 1 3 5 5
##  $ LandSlope      : Factor w/ 3 levels "Gtl","Mod","Sev": 1 1 1 1 1 1 1 1
##  $ Neighborhood   : Factor w/ 25 levels "Blmngtn","Blueste",...: 6 25 6 7 14
##  $ Condition1     : Factor w/ 9 levels "Artery","Feedr",...: 3 2 3 3 3 3 3 5
##  $ Condition2     : Factor w/ 8 levels "Artery","Feedr",...: 3 3 3 3 3 3 3 3
##  $ BldgType       : Factor w/ 5 levels "1Fam","2fmCon",...: 1 1 1 1 1 1 1 1
##  $ HouseStyle     : Factor w/ 8 levels "1.5Fin","1.5Unf",...: 6 3 6 6 6 1 3 6
##  $ OverallQual    : int   7 6 7 7 8 5 8 7 7 5 ...
##  $ OverallCond    : int   5 8 5 5 5 5 5 6 5 6 ...
##  $ YearBuilt      : int   2003 1976 2001 1915 2000 1993 2004 1973 1931 1939
##  $ YearRemodAdd   : int   2003 1976 2002 1970 2000 1995 2005 1973 1950 1950
##  $ RoofStyle      : Factor w/ 6 levels "Flat","Gable",...: 2 2 2 2 2 2 2 2
##  $ RoofMatl       : Factor w/ 8 levels "ClyTile","CompShg",...: 2 2 2 2 2 2 2 2
##  $ Exterior1st    : Factor w/ 15 levels "AsbShng","AsphShn",...: 13 9 13 14
```

```

13 13 13 7 4 9 ...
## $ Exterior2nd : Factor w/ 16 levels "AsbShng","AsphShn",...: 14 9 14 16
14 14 14 7 16 9 ...
## $ MasVnrType : Factor w/ 4 levels "BrkCmn","BrkFace",...: 2 3 2 3 2 3 4
4 3 3 ...
## $ MasVnrArea : int 196 0 162 0 350 0 186 240 0 0 ...
## $ ExterQual : Factor w/ 4 levels "Ex","Fa","Gd",...: 3 4 3 4 3 4 3 4 4
4 ...
## $ ExterCond : Factor w/ 5 levels "Ex","Fa","Gd",...: 5 5 5 5 5 5 5 5 5
5 ...
## $ Foundation : Factor w/ 6 levels "BrkTil","CBlock",...: 3 2 3 1 3 6 3 2
1 1 ...
## $ BsmtQual : Factor w/ 4 levels "Ex","Fa","Gd",...: 3 3 3 4 3 3 1 3 4
4 ...
## $ BsmtCond : Factor w/ 4 levels "Fa","Gd","Po",...: 4 4 4 2 4 4 4 4 4
4 ...
## $ BsmtExposure : Factor w/ 4 levels "Av","Gd","Mn",...: 4 2 3 4 1 4 1 3 4
4 ...
## $ BsmtFinType1 : Factor w/ 6 levels "ALQ","BLQ","GLQ",...: 3 1 3 1 3 3 3 1
6 3 ...
## $ BsmtFinSF1 : int 706 978 486 216 655 732 1369 859 0 851 ...
## $ BsmtFinType2 : Factor w/ 6 levels "ALQ","BLQ","GLQ",...: 6 6 6 6 6 6 6 2
6 6 ...
## $ BsmtFinSF2 : int 0 0 0 0 0 0 0 32 0 0 ...
## $ BsmtUnfSF : int 150 284 434 540 490 64 317 216 952 140 ...
## $ TotalBsmtSF : int 856 1262 920 756 1145 796 1686 1107 952 991 ...
## $ Heating : Factor w/ 6 levels "Floor","GasA",...: 2 2 2 2 2 2 2 2 2
2 ...
## $ HeatingQC : Factor w/ 5 levels "Ex","Fa","Gd",...: 1 1 1 3 1 1 1 1 3
1 ...
## $ CentralAir : Factor w/ 2 levels "N","Y": 2 2 2 2 2 2 2 2 2 2 ...
## $ Electrical : Factor w/ 5 levels "FuseA","FuseF",...: 5 5 5 5 5 5 5 5 2
5 ...
## $ X1stFlrSF : int 856 1262 920 961 1145 796 1694 1107 1022 1077 ...
## $ X2ndFlrSF : int 854 0 866 756 1053 566 0 983 752 0 ...
## $ LowQualFinSF : int 0 0 0 0 0 0 0 0 0 0 ...
## $ GrLivArea : int 1710 1262 1786 1717 2198 1362 1694 2090 1774 1077
...
## $ BsmtFullBath : int 1 0 1 1 1 1 1 1 0 1 ...
## $ BsmtHalfBath : int 0 1 0 0 0 0 0 0 0 0 ...
## $ FullBath : int 2 2 2 1 2 1 2 2 2 1 ...
## $ HalfBath : int 1 0 1 0 1 1 0 1 0 0 ...
## $ BedroomAbvGr : int 3 3 3 3 4 1 3 3 2 2 ...
## $ KitchenAbvGr : int 1 1 1 1 1 1 1 1 2 2 ...
## $ KitchenQual : Factor w/ 4 levels "Ex","Fa","Gd",...: 3 4 3 3 3 4 3 4 4
4 ...
## $ TotRmsAbvGrd : int 8 6 6 7 9 5 7 7 8 5 ...
## $ Functional : Factor w/ 7 levels "Maj1","Maj2",...: 7 7 7 7 7 7 7 7 3 7
...
## $ Fireplaces : int 0 1 1 1 1 0 1 2 2 2 ...

```

```

## $ FireplaceQu : Factor w/ 5 levels "Ex","Fa","Gd",...: NA 5 5 3 5 NA 3 5
5 5 ...
## $ GarageType : Factor w/ 6 levels "2Types","Attchd",...: 2 2 2 6 2 2 2 2
6 2 ...
## $ GarageYrBlt : int 2003 1976 2001 1998 2000 1993 2004 1973 1931 1939
...
## $ GarageFinish : Factor w/ 3 levels "Fin","RFn","Unf": 2 2 2 3 2 3 2 2 3
2 ...
## $ GarageCars : int 2 2 2 3 3 2 2 2 2 1 ...
## $ GarageArea : int 548 460 608 642 836 480 636 484 468 205 ...
## $ GarageQual : Factor w/ 5 levels "Ex","Fa","Gd",...: 5 5 5 5 5 5 5 5 2
3 ...
## $ GarageCond : Factor w/ 5 levels "Ex","Fa","Gd",...: 5 5 5 5 5 5 5 5 5
5 ...
## $ PavedDrive : Factor w/ 3 levels "N","P","Y": 3 3 3 3 3 3 3 3 3 3 ...
## $ WoodDeckSF : int 0 298 0 0 192 40 255 235 90 0 ...
## $ OpenPorchSF : int 61 0 42 35 84 30 57 204 0 4 ...
## $ EnclosedPorch: int 0 0 0 272 0 0 0 228 205 0 ...
## $ X3SsnPorch : int 0 0 0 0 0 320 0 0 0 0 ...
## $ ScreenPorch : int 0 0 0 0 0 0 0 0 0 0 ...
## $ PoolArea : int 0 0 0 0 0 0 0 0 0 0 ...
## $ PoolQC : Factor w/ 3 levels "Ex","Fa","Gd": NA NA NA NA NA NA NA
NA NA NA ...
## $ Fence : Factor w/ 4 levels "GdPrv","GdWo",...: NA NA NA NA NA 3
NA NA NA NA ...
## $ MiscFeature : Factor w/ 4 levels "Gar2","Othr",...: NA NA NA NA NA 3 NA
3 NA NA ...
## $ MiscVal : int 0 0 0 0 0 700 0 350 0 0 ...
## $ MoSold : int 2 5 9 2 12 10 8 11 4 1 ...
## $ YrSold : int 2008 2007 2008 2006 2008 2009 2007 2009 2008 2008
...
## $ SaleType : Factor w/ 9 levels "COD","Con","ConLD",...: 9 9 9 9 9 9 9
9 9 9 ...
## $ SaleCondition: Factor w/ 6 levels "Abnorml","AdjLand",...: 5 5 5 1 5 5 5
5 1 5 ...
## $ SalePrice : int 208500 181500 223500 140000 250000 143000 307000
200000 129900 118000 ...

```

We can see that our data consists of only integer and Factor columns, meaning that they are either whole numbers or some text description of the house.

```
summary(train_data)
```

```

##      Id      MSSubClass      MSZoning      LotFrontage
## Min.   : 1.0   Min.   : 20.0   C (all): 10   Min.   : 21.00
## 1st Qu.: 365.8 1st Qu.: 20.0   FV    : 65   1st Qu.: 59.00
## Median : 730.5 Median : 50.0   RH    : 16   Median : 69.00
## Mean   : 730.5 Mean   : 56.9   RL    :1151  Mean   : 70.05
## 3rd Qu.:1095.2 3rd Qu.: 70.0   RM    : 218  3rd Qu.: 80.00
## Max.   :1460.0 Max.   :190.0           Max.   :313.00

```

```

##                                     NA's :259
##      LotArea      Street      Alley      LotShape      LandContour      Utilities
## Min. : 1300      Grvl: 6      Grvl: 50      IR1:484      Bnk: 63
AllPub:1459
## 1st Qu.: 7554      Pave:1454      Pave: 41      IR2: 41      HLS: 50      NoSeWa:
1
## Median : 9478                                     NA's:1369      IR3: 10      Low: 36
## Mean : 10517                                     Reg:925      Lvl:1311
## 3rd Qu.: 11602
## Max. :215245
##
##      LotConfig      LandSlope      Neighborhood      Condition1      Condition2
## Corner : 263      Gtl:1382      NAmes :225      Norm :1260      Norm :1445
## CulDSac: 94      Mod: 65      CollgCr:150      Feedr : 81      Feedr : 6
## FR2 : 47      Sev: 13      OldTown:113      Artery : 48      Artery : 2
## FR3 : 4                                     Edwards:100      RRAn : 26      PosN : 2
## Inside :1052      Somerst: 86      PosN : 19      RRNn : 2
##                                     Gilbert: 79      RRAe : 11      PosA : 1
##                                     (Other):707      (Other): 15      (Other): 2
##      BldgType      HouseStyle      OverallQual      OverallCond      YearBuilt
## 1Fam :1220      1Story :726      Min. : 1.000      Min. :1.000      Min. :1872
## 2fmCon: 31      2Story :445      1st Qu.: 5.000      1st Qu.:5.000      1st Qu.:1954
## Duplex: 52      1.5Fin :154      Median : 6.000      Median :5.000      Median :1973
## Twnhs : 43      SLvl : 65      Mean : 6.099      Mean :5.575      Mean :1971
## TwnhsE: 114      SFoyer : 37      3rd Qu.: 7.000      3rd Qu.:6.000      3rd Qu.:2000
##                                     1.5Unf : 14      Max. :10.000      Max. :9.000      Max. :2010
##                                     (Other): 19
##      YearRemodAdd      RoofStyle      RoofMatl      Exterior1st      Exterior2nd
## Min. :1950      Flat : 13      CompShg:1434      VinylSd:515      VinylSd:504
## 1st Qu.:1967      Gable :1141      Tar&Grv: 11      HdBoard:222      MetalSd:214
## Median :1994      Gambrel: 11      WdShngl: 6      MetalSd:220      HdBoard:207
## Mean :1985      Hip : 286      WdShake: 5      Wd Sdng:206      Wd Sdng:197
## 3rd Qu.:2004      Mansard: 7      ClyTile: 1      Plywood:108      Plywood:142
## Max. :2010      Shed : 2      Membran: 1      CemntBd: 61      CmentBd: 60
##                                     (Other): 2      (Other):128      (Other):136
##      MasVnrType      MasVnrArea      ExterQual      ExterCond      Foundation      BsmtQual
## BrkCmn : 15      Min. : 0.0      Ex: 52      Ex: 3      BrkTil:146      Ex :121
## BrkFace:445      1st Qu.: 0.0      Fa: 14      Fa: 28      CBlock:634      Fa : 35
## None :864      Median : 0.0      Gd:488      Gd: 146      PConc :647      Gd :618
## Stone :128      Mean : 103.7      TA:906      Po: 1      Slab : 24      TA :649
## NA's : 8      3rd Qu.: 166.0      TA:1282      Stone : 6      NA's: 37
##                                     Max. :1600.0      Wood : 3
##                                     NA's :8
##      BsmtCond      BsmtExposure      BsmtFinType1      BsmtFinSF1      BsmtFinType2
## Fa : 45      Av :221      ALQ :220      Min. : 0.0      ALQ : 19
## Gd : 65      Gd :134      BLQ :148      1st Qu.: 0.0      BLQ : 33
## Po : 2      Mn :114      GLQ :418      Median : 383.5      GLQ : 14
## TA :1311      No :953      LwQ : 74      Mean : 443.6      LwQ : 46
## NA's: 37      NA's: 38      Rec :133      3rd Qu.: 712.2      Rec : 54
##                                     Unf :430      Max. :5644.0      Unf :1256

```

```

##                                     NA's: 37                                     NA's: 38
##      BsmtFinSF2      BsmtUnfSF      TotalBsmtSF      Heating
HeatingQC
## Min.   : 0.00   Min.   : 0.0   Min.   : 0.0   Floor:   1   Ex:741
## 1st Qu.: 0.00   1st Qu.: 223.0   1st Qu.: 795.8   GasA :1428   Fa: 49
## Median : 0.00   Median : 477.5   Median : 991.5   GasW : 18   Gd:241
## Mean   : 46.55   Mean   : 567.2   Mean   :1057.4   Grav : 7    Po: 1
## 3rd Qu.: 0.00   3rd Qu.: 808.0   3rd Qu.:1298.2   OthW : 2    TA:428
## Max.   :1474.00   Max.   :2336.0   Max.   :6110.0   Wall : 4
##
##      CentralAir Electrical      X1stFlrSF      X2ndFlrSF      LowQualFinSF
## N: 95      FuseA: 94   Min.   : 334   Min.   : 0    Min.   : 0.000
## Y:1365      FuseF: 27   1st Qu.: 882   1st Qu.: 0    1st Qu.: 0.000
##              FuseP: 3   Median :1087   Median : 0    Median : 0.000
##              Mix : 1   Mean   :1163   Mean   : 347   Mean   : 5.845
##              SBrkr:1334 3rd Qu.:1391   3rd Qu.: 728   3rd Qu.: 0.000
##              NA's : 1   Max.   :4692   Max.   :2065   Max.   :572.000
##
##      GrLivArea      BsmtFullBath      BsmtHalfBath      FullBath
## Min.   : 334   Min.   :0.0000   Min.   :0.0000   Min.   :0.000
## 1st Qu.:1130   1st Qu.:0.0000   1st Qu.:0.0000   1st Qu.:1.000
## Median :1464   Median :0.0000   Median :0.0000   Median :2.000
## Mean   :1515   Mean   :0.4253   Mean   :0.05753   Mean   :1.565
## 3rd Qu.:1777   3rd Qu.:1.0000   3rd Qu.:0.0000   3rd Qu.:2.000
## Max.   :5642   Max.   :3.0000   Max.   :2.0000   Max.   :3.000
##
##      HalfBath      BedroomAbvGr      KitchenAbvGr      KitchenQual
TotRmsAbvGrd
## Min.   :0.0000   Min.   :0.000   Min.   :0.000   Ex:100   Min.   :
2.000
## 1st Qu.:0.0000   1st Qu.:2.000   1st Qu.:1.000   Fa: 39   1st Qu.:
5.000
## Median :0.0000   Median :3.000   Median :1.000   Gd:586   Median :
6.000
## Mean   :0.3829   Mean   :2.866   Mean   :1.047   TA:735   Mean   :
6.518
## 3rd Qu.:1.0000   3rd Qu.:3.000   3rd Qu.:1.000   3rd Qu.:
7.000
## Max.   :2.0000   Max.   :8.000   Max.   :3.000   Max.
:14.000
##
##      Functional      Fireplaces      FireplaceQu      GarageType      GarageYrBlt
## Maj1: 14   Min.   :0.000   Ex : 24   2Types : 6   Min.   :1900
## Maj2: 5    1st Qu.:0.000   Fa : 33   Attchd :870   1st Qu.:1961
## Min1: 31   Median :1.000   Gd :380   Basement: 19   Median :1980
## Min2: 34   Mean   :0.613   Po : 20   BuiltIn: 88   Mean   :1979
## Mod : 15   3rd Qu.:1.000   TA :313   CarPort: 9    3rd Qu.:2002
## Sev : 1    Max.   :3.000   NA's:690   Detchd :387   Max.   :2010
## Typ :1360                                     NA's : 81   NA's :81
##      GarageFinish      GarageCars      GarageArea      GarageQual      GarageCond

```



```

## Fin :352      Min.    :0.000  Min.    :  0.0  Ex   :   3  Ex   :   2
## RFn :422      1st Qu.:1.000  1st Qu.: 334.5  Fa   :  48  Fa   :  35
## Unf :605      Median :2.000  Median : 480.0  Gd   :  14  Gd   :   9
## NA's: 81      Mean    :1.767  Mean    : 473.0  Po   :   3  Po   :   7
##              3rd Qu.:2.000  3rd Qu.: 576.0  TA   :1311  TA   :1326
##              Max.    :4.000  Max.    :1418.0  NA's:  81  NA's:  81
##
## PavedDrive    WoodDeckSF      OpenPorchSF      EnclosedPorch
X3SsnPorch
## N:  90      Min.    :  0.00  Min.    :  0.00  Min.    :  0.00  Min.    :
0.00
## P:  30      1st Qu.:  0.00  1st Qu.:  0.00  1st Qu.:  0.00  1st Qu.:
0.00
## Y:1340      Median :  0.00  Median : 25.00  Median :  0.00  Median :
0.00
##              Mean    : 94.24  Mean    : 46.66  Mean    : 21.95  Mean    :
3.41
##              3rd Qu.:168.00  3rd Qu.: 68.00  3rd Qu.:  0.00  3rd Qu.:
0.00
##              Max.    :857.00  Max.    :547.00  Max.    :552.00  Max.
:508.00
##
## ScreenPorch    PoolArea      PoolQC      Fence      MiscFeature
## Min.    :  0.00  Min.    :  0.000  Ex   :   2  GdPrv:  59  Gar2:   2
## 1st Qu.:  0.00  1st Qu.:  0.000  Fa   :   2  GdWo :  54  Othr:   2
## Median :  0.00  Median :  0.000  Gd   :   3  MnPrv: 157  Shed:  49
## Mean    : 15.06  Mean    :  2.759  NA's:1453  MnWw :  11  TenC:   1
## 3rd Qu.:  0.00  3rd Qu.:  0.000  NA's :1179  NA's:1406
## Max.    :480.00  Max.    :738.000
##
## MiscVal      MoSold      YrSold      SaleType
## Min.    :  0.00  Min.    :  1.000  Min.    :2006  WD      :1267
## 1st Qu.:  0.00  1st Qu.:  5.000  1st Qu.:2007  New     : 122
## Median :  0.00  Median :  6.000  Median :2008  COD     :  43
## Mean    :  43.49  Mean    :  6.322  Mean    :2008  ConLD   :   9
## 3rd Qu.:  0.00  3rd Qu.:  8.000  3rd Qu.:2009  ConLI   :   5
## Max.    :15500.00  Max.    :12.000  Max.    :2010  ConLw   :   5
##              (Other):   9
## SaleCondition  SalePrice
## Abnorml: 101  Min.    : 34900
## AdjLand:   4  1st Qu.:129975
## Alloca :  12  Median :163000
## Family  :  20  Mean    :180921
## Normal :1198  3rd Qu.:214000
## Partial: 125  Max.    :755000
##

```

From this command, we can see the summary from every column. More importantly, the summary statistics of the numerical columns. We can see that the average price of a home

is 180,921. We can also see that an average home was built in 1971 with around 2.8 rooms above ground.

Data Cleaning

As seen above, this dataset has rows with missing values, let's check how many missing values there actually are.

```
NA_values = data.frame(NA_value=colSums(is.na(train_data)))
NA_values
```

| ## | NA_value |
|-----------------|----------|
| ## Id | 0 |
| ## MSSubClass | 0 |
| ## MSZoning | 0 |
| ## LotFrontage | 259 |
| ## LotArea | 0 |
| ## Street | 0 |
| ## Alley | 1369 |
| ## LotShape | 0 |
| ## LandContour | 0 |
| ## Utilities | 0 |
| ## LotConfig | 0 |
| ## LandSlope | 0 |
| ## Neighborhood | 0 |
| ## Condition1 | 0 |
| ## Condition2 | 0 |
| ## BldgType | 0 |
| ## HouseStyle | 0 |
| ## OverallQual | 0 |
| ## OverallCond | 0 |
| ## YearBuilt | 0 |
| ## YearRemodAdd | 0 |
| ## RoofStyle | 0 |
| ## RoofMatl | 0 |
| ## Exterior1st | 0 |
| ## Exterior2nd | 0 |
| ## MasVnrType | 8 |
| ## MasVnrArea | 8 |
| ## ExterQual | 0 |
| ## ExterCond | 0 |
| ## Foundation | 0 |
| ## BsmtQual | 37 |
| ## BsmtCond | 37 |
| ## BsmtExposure | 38 |
| ## BsmtFinType1 | 37 |
| ## BsmtFinSF1 | 0 |
| ## BsmtFinType2 | 38 |
| ## BsmtFinSF2 | 0 |

```
## BsmtUnfSF      0
## TotalBsmtSF    0
## Heating        0
## HeatingQC      0
## CentralAir     0
## Electrical     1
## X1stFlrSF      0
## X2ndFlrSF      0
## LowQualFinSF   0
## GrLivArea      0
## BsmtFullBath   0
## BsmtHalfBath   0
## FullBath       0
## HalfBath       0
## BedroomAbvGr   0
## KitchenAbvGr   0
## KitchenQual    0
## TotRmsAbvGrd   0
## Functional     0
## Fireplaces     0
## FireplaceQu    690
## GarageType     81
## GarageYrBlt    81
## GarageFinish   81
## GarageCars     0
## GarageArea     0
## GarageQual     81
## GarageCond     81
## PavedDrive     0
## WoodDeckSF     0
## OpenPorchSF    0
## EnclosedPorch  0
## X3SsnPorch     0
## ScreenPorch    0
## PoolArea       0
## PoolQC        1453
## Fence         1179
## MiscFeature    1406
## MiscVal       0
## MoSold        0
## YrSold        0
## SaleType      0
## SaleCondition 0
## SalePrice     0
```

Let's drop any columns with missing values in both the training and test set, as well as the ID column as it is not useful.

```
# Getting rid of training and testing columns with missing values
train_data = subset(train_data, select = -
```

```
c(Id, LotFrontage, Alley, MasVnrType, MasVnrArea, BsmtQual, BsmtCond, BsmtExposure, B
smtFinType1, BsmtFinType2, Electrical, FireplaceQu, GarageType, GarageYrBlt, Garage
Finish, GarageQual, GarageCond, PoolQC, Fence, MiscFeature))
```

```
test_data = subset(test_dataa, select = -
c(Id, LotFrontage, Alley, MasVnrType, MasVnrArea, BsmtQual, BsmtCond, BsmtExposure, B
smtFinType1, BsmtFinType2, Electrical, FireplaceQu, GarageType, GarageYrBlt, Garage
Finish, GarageQual, GarageCond, PoolQC, Fence, MiscFeature))
```

Exploratory Data Analysis on Training Set

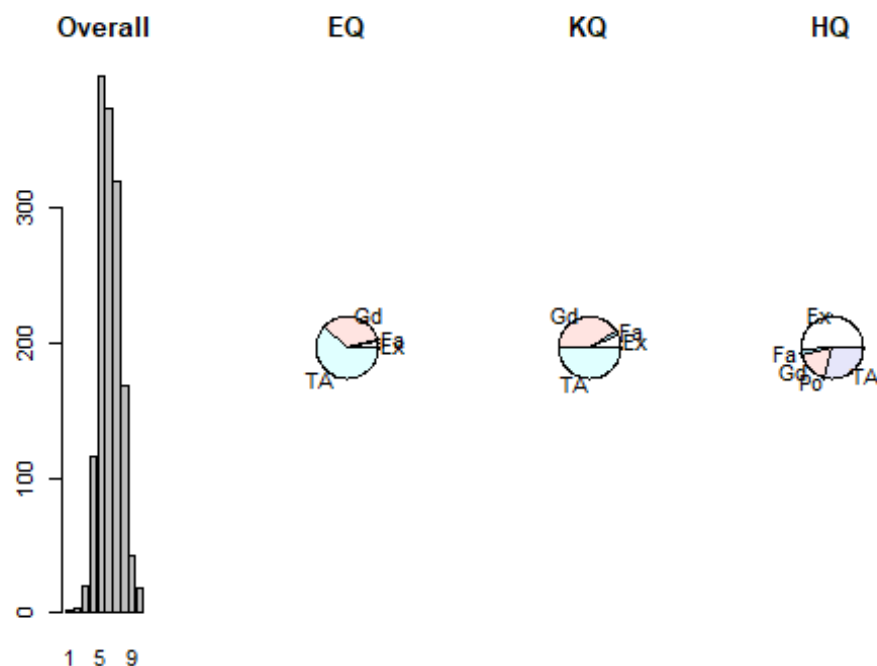
Pie Plot

```
par(mfrow=c(1,4))
barplot(table(train_data$OverallQual), main="Overall")
```

```
pie(table(train_data$ExterQual), labels =
names(table(train_data$ExterQual)), main="EQ")
```

```
pie(table(train_data$KitchenQual), labels =
names(table(train_data$KitchenQual)), main="KQ")
```

```
pie(table(train_data$HeatingQC), labels =
names(table(train_data$HeatingQC)), main="HQ")
```




As we can see, most houses have quality of 5-6 which are average. This coincides with our kitchen and heating

quality as most are “TA” or average. Surprisingly, most houses have excellent exterior quality. This may be due to how much the kitchen and heating get used over time so it wears down.

Correlation Matrix Plot

```
numerical_data <- train_data %>% dplyr::select(where(is.numeric))
cor_data=data.frame(numerical_data)
correlation = cor(cor_data)
par(mfrow=c(1,1))
corrplot(correlation,method="color", tl.cex = 0.7)
```



In our correlation plot of only our numerical columns, we can see some interesting findings. Importantly, the SalePrice column has some very strong positive correlations with the overall quality of the house (OverallQual) and the above ground living area (GrLivArea), with also some strong correlations with most columns. There are columns with not much correlation with SalePrice as we can see, the building class (MSSubClass) has almost no correlation since the color is white. Most columns after the Enclosed Porch column have no correlation to SalePrice. There are not many negative correlations to SalePrice aswell.

Strongly Correlated Columns Boxplot

```
par(mfrow=c(3,1))
par(plt = c(0.1, 0.9, 0.3, 1))

boxplot(SalePrice~OverallQual, data=train_data, col="lightblue",
border="cadetblue4",
main="SalePrice and Overall Quality",
```


```

        xlab="Overall Quality", ylab="Price")

boxplot(SalePrice~TotRmsAbvGrd, data=train_data, col="lightblue",
border="cadetblue4",
        main="SalePrice and #Above Ground Rooms",
        xlab="Total Rooms Above Ground", ylab="Price")

boxplot(SalePrice~FullBath, data=train_data, col="lightblue",
border="cadetblue4",
        main="SalePrice and #Full Bathrooms",
        xlab="Full Bathrooms", ylab="Price")

```



We can see the general trend in these boxplots. When the overall quality increases on average, so does the price of a house as expected. The total above ground rooms, there seems to be a positive correlation but the price goes down when there are 12-14 rooms for some reason. This could be due to different factors of each house individually which cause a decrease in the overall median price. When there is 1 full bathroom, it seems to have the same median price as a house with 0 full baths, which is quite unexpected.

```

print(sum(train_data$FullBath == 0))

## [1] 9

print(sum(train_data$FullBath == 1))


## [1] 650

```

As shown above, the it is hard to compare 0 Full bathroom houses with 1 Full Bathroom houses since there are not many houses with 0 full bathrooms in the dataset, hence why they have similar median prices.

Histogram of SalePrice to see Price distribution


```
options(scipen = 999)
hist(train_data$SalePrice, col="lightblue")
```



As we can see, most of the houses in this dataset are \$100,000 - \$200,000.

SalePrice Boxplot to see outliers

```
boxplot(train_data$SalePrice, col="lightblue", border = "cadetblue4")
```



As we can see, the outliers are above around \$350,000, while most houses range between around \$150,000 to \$220,000. We notice that the number of outliers may not be significant enough in comparison to the entire dataset to remove.

Data Altering

There are a lot of columns that are not correlated to the price of a house. Since we are using a linear classifier, we should see the effects of a raw dataset with all of our columns, and another dataset with reduced columns that have no correlation ones removed. We should also check if the number of outliers are significant to our entire dataset.

```
message("Total Number of Training Data: ", nrow(train_data))  
## Total Number of Training Data: 1460  
message("Number of outliers: ", sum(train_data$SalePrice > 350000))  
## Number of outliers: 54
```

As we can see, there are clearly not enough significant outliers that can heavily alter our dataset, so it is not worth removing them.

Now, Let's create two linear models, one which has every column in the dataset, while the other excludes columns that have low correlation with SalePrice


```
numerical_data <- train_data %>% dplyr::select(where(is.numeric))
cor_data=data.frame(numerical_data)
print(cor(cor_data$SalePrice, cor_data))

##      MSSubClass   LotArea OverallQual OverallCond YearBuilt YearRemodAdd
## [1,] -0.08428414 0.2638434   0.7909816 -0.07785589 0.5228973   0.507101
##      BsmtFinSF1  BsmtFinSF2 BsmtUnfSF TotalBsmtSF X1stFlrSF X2ndFlrSF
## [1,]  0.3864198 -0.01137812 0.2144791   0.6135806 0.6058522 0.3193338
##      LowQualFinSF GrLivArea BsmtFullBath BsmtHalfBath FullBath HalfBath
## [1,] -0.02560613 0.7086245   0.2271222 -0.01684415 0.5606638 0.2841077
##      BedroomAbvGr KitchenAbvGr TotRmsAbvGrd Fireplaces GarageCars
GarageArea
## [1,]  0.1682132  -0.1359074    0.5337232  0.4669288  0.6404092
0.6234314
##      WoodDeckSF OpenPorchSF EnclosedPorch X3SsnPorch ScreenPorch
PoolArea
## [1,]  0.3244134  0.3158562    -0.128578 0.04458367  0.1114466
0.09240355
##      MiscVal      MoSold      YrSold SalePrice
## [1,] -0.02118958 0.04643225 -0.02892259      1
```

We should try a model with columns of correlation above absolute value of 0.3. This way, we do not lose too many features while also retaining the most important factors of house price.

Validation Set

Let's also create a validation set since the test set does not have labels in which we can check our metrics

```
set.seed(42)
sample = sample.split(cor_data, SplitRatio = 0.9)
train.data = subset(cor_data, sample==TRUE)
val.data = subset(cor_data, sample==FALSE)

numerical_t <- test_data %>% dplyr::select(where(is.numeric))
cor_dat=data.frame(numerical_t)

cor_dat[is.na(cor_dat)] <- 0
test.data <- model.matrix(~.,cor_dat)[,-1]
```

Data Modelling

Linear Regression 1

Firstly, Lets create a general Linear model. One with every column and another with only columns with high correlation.

```

model_orig = lm(SalePrice ~ MSSubClass+LotArea+OverallQual+OverallCond
+YearBuilt+BsmFinSF1+YearRemodAdd+BsmFinSF2+BsmUnfSF+TotalBsmSF+X1stFlrSF
+X2ndFlrSF+LowQualFinSF +GrLivArea +
BsmFullBath+BsmHalfBath+FullBath+HalfBath+BedroomAbvGr+KitchenAbvGr
+TotRmsAbvGrd +Fireplaces +GarageCars +GarageArea+WoodDeckSF +OpenPorchSF
+EnclosedPorch + X3SsnPorch +ScreenPorch+PoolArea+MiscVal + MoSold+
YrSold,data = train.data)

model_reduced = lm(SalePrice ~ OverallQual
+YearBuilt+BsmFinSF1+YearRemodAdd+TotalBsmSF+X1stFlrSF +X2ndFlrSF
+GrLivArea +FullBath+TotRmsAbvGrd +Fireplaces +GarageCars
+GarageArea+WoodDeckSF +OpenPorchSF,data = train.data)

```

Lasso Regression

Let's create a Lasso regression model. We will compare all these models in the next part.

```

library(glmnet)

## Warning: package 'glmnet' was built under R version 4.2.3
## Loading required package: Matrix
## Warning: package 'Matrix' was built under R version 4.2.3
##
## Attaching package: 'Matrix'
## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack
## Loaded glmnet 4.1-8

set.seed(42)
x <- model.matrix(SalePrice~., train.data)[,-1]
y <- train.data$SalePrice

x_val <- model.matrix(SalePrice~., val.data)[,-1]
y_val <- val.data$SalePrice

cv <- cv.glmnet(x,y,alpha=1)

# Fit the model on training data using Lowest Lambda
modelLass <- glmnet(x,y,alpha=1, lambda=cv$lambda.min)
coef(modelLass)

## 34 x 1 sparse Matrix of class "dgCMatrix"
##                                     s0
## (Intercept)      -847338.2958875
## MSSubClass        -88.0074685
## LotArea           0.1948199

```

```
## OverallQual      21075.3152618
## OverallCond      .
## YearBuilt        161.7608713
## YearRemodAdd     237.6645869
## BsmtFinSF1       14.1611180
## BsmtFinSF2       .
## BsmtUnfSF        .
## TotalBsmtSF      10.1828146
## X1stFlrSF        5.0687187
## X2ndFlrSF        .
## LowQualFinSF     .
## GrLivArea        42.9509141
## BsmtFullBath     2936.9278669
## BsmtHalfBath     .
## FullBath         .
## HalfBath         .
## BedroomAbvGr     .
## KitchenAbvGr     -2264.6577129
## TotRmsAbvGrd     .
## Fireplaces       3936.6657142
## GarageCars       11023.7857163
## GarageArea       2.6579724
## WoodDeckSF       19.6422321
## OpenPorchSF      .
## EnclosedPorch    .
## X3SsnPorch       .
## ScreenPorch      .
## PoolArea         .
## MiscVal          .
## MoSold           .
## YrSold           .
```

Elastic Net Regression

```
set.seed(42)
modelNet <- train(SalePrice ~., data=train.data, method="glmnet",
trControl=trainControl("cv",number=10), tuneLength = 10)

modelNet

## glmnet
##
## 1288 samples
## 33 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1159, 1159, 1160, 1159, 1159, ...
## Resampling results across tuning parameters:
##
##  alpha  lambda      RMSE      Rsquared    MAE
```

| | | | | | |
|----|-----|------------|----------|-----------|----------|
| ## | 0.1 | 67.1298 | 36112.99 | 0.8006632 | 21927.71 |
| ## | 0.1 | 155.0785 | 36112.99 | 0.8006632 | 21927.71 |
| ## | 0.1 | 358.2515 | 36112.99 | 0.8006632 | 21927.71 |
| ## | 0.1 | 827.6075 | 36060.86 | 0.8011534 | 21880.06 |
| ## | 0.1 | 1911.8807 | 35985.81 | 0.8019241 | 21783.73 |
| ## | 0.1 | 4416.6923 | 35891.08 | 0.8030804 | 21615.30 |
| ## | 0.1 | 10203.1320 | 35999.80 | 0.8032724 | 21518.41 |
| ## | 0.1 | 23570.5583 | 37078.83 | 0.7976901 | 22009.34 |
| ## | 0.1 | 54451.0468 | 40018.02 | 0.7892309 | 23774.93 |
| ## | 0.2 | 67.1298 | 36116.25 | 0.8006069 | 21932.01 |
| ## | 0.2 | 155.0785 | 36116.25 | 0.8006069 | 21932.01 |
| ## | 0.2 | 358.2515 | 36096.36 | 0.8007989 | 21913.76 |
| ## | 0.2 | 827.6075 | 36036.85 | 0.8013884 | 21844.75 |
| ## | 0.2 | 1911.8807 | 35968.36 | 0.8021008 | 21733.75 |
| ## | 0.2 | 4416.6923 | 35941.70 | 0.8027885 | 21617.55 |
| ## | 0.2 | 10203.1320 | 36578.11 | 0.7987037 | 21872.95 |
| ## | 0.2 | 23570.5583 | 38298.98 | 0.7916543 | 22774.16 |
| ## | 0.2 | 54451.0468 | 43271.31 | 0.7819889 | 26435.92 |
| ## | 0.3 | 67.1298 | 36117.51 | 0.8005950 | 21933.63 |
| ## | 0.3 | 155.0785 | 36117.51 | 0.8005950 | 21933.63 |
| ## | 0.3 | 358.2515 | 36080.94 | 0.8009492 | 21895.77 |
| ## | 0.3 | 827.6075 | 36025.15 | 0.8014950 | 21817.96 |
| ## | 0.3 | 1911.8807 | 35956.46 | 0.8022339 | 21717.65 |
| ## | 0.3 | 4416.6923 | 36127.03 | 0.8012456 | 21726.58 |
| ## | 0.3 | 10203.1320 | 37044.19 | 0.7955403 | 22205.69 |
| ## | 0.3 | 23570.5583 | 39468.24 | 0.7870941 | 23588.58 |
| ## | 0.3 | 54451.0468 | 46989.09 | 0.7725753 | 29905.05 |
| ## | 0.4 | 67.1298 | 36119.96 | 0.8005900 | 21931.95 |
| ## | 0.4 | 155.0785 | 36116.32 | 0.8006217 | 21928.87 |
| ## | 0.4 | 358.2515 | 36067.88 | 0.8010770 | 21877.39 |
| ## | 0.4 | 827.6075 | 36014.17 | 0.8015980 | 21795.33 |
| ## | 0.4 | 1911.8807 | 35956.20 | 0.8022663 | 21712.36 |
| ## | 0.4 | 4416.6923 | 36379.86 | 0.7990023 | 21898.63 |
| ## | 0.4 | 10203.1320 | 37549.01 | 0.7921979 | 22493.84 |
| ## | 0.4 | 23570.5583 | 40712.30 | 0.7820659 | 24553.03 |
| ## | 0.4 | 54451.0468 | 50695.25 | 0.7643866 | 33332.98 |
| ## | 0.5 | 67.1298 | 36119.02 | 0.8006149 | 21930.55 |
| ## | 0.5 | 155.0785 | 36105.37 | 0.8007176 | 21919.55 |
| ## | 0.5 | 358.2515 | 36055.52 | 0.8011944 | 21859.73 |
| ## | 0.5 | 827.6075 | 36000.83 | 0.8017118 | 21778.17 |
| ## | 0.5 | 1911.8807 | 35998.79 | 0.8018873 | 21731.66 |
| ## | 0.5 | 4416.6923 | 36686.91 | 0.7961497 | 22126.99 |
| ## | 0.5 | 10203.1320 | 38057.10 | 0.7887529 | 22852.08 |
| ## | 0.5 | 23570.5583 | 42025.36 | 0.7759593 | 25639.02 |
| ## | 0.5 | 54451.0468 | 54404.90 | 0.7532590 | 36548.39 |
| ## | 0.6 | 67.1298 | 36116.20 | 0.8006290 | 21929.53 |
| ## | 0.6 | 155.0785 | 36095.63 | 0.8008036 | 21909.99 |
| ## | 0.6 | 358.2515 | 36049.00 | 0.8012481 | 21845.63 |
| ## | 0.6 | 827.6075 | 35985.35 | 0.8018504 | 21760.32 |
| ## | 0.6 | 1911.8807 | 36078.97 | 0.8011955 | 21766.76 |

```
## 0.6      4416.6923  36899.12  0.7944217  22267.97
## 0.6     10203.1320  38577.09  0.7850389  23250.00
## 0.6     23570.5583  43503.89  0.7665794  26922.41
## 0.6     54451.0468  58460.50  0.7259307  39979.16
## 0.7         67.1298  36117.31  0.8006142  21930.26
## 0.7        155.0785  36089.05  0.8008679  21901.09
## 0.7        358.2515  36044.81  0.8012809  21834.13
## 0.7        827.6075  35974.40  0.8019493  21749.27
## 0.7       1911.8807  36179.50  0.8002798  21814.85
## 0.7       4416.6923  37078.60  0.7930379  22364.24
## 0.7      10203.1320  39072.37  0.7815631  23655.41
## 0.7      23570.5583  44963.10  0.7569792  28240.16
## 0.7      54451.0468  62166.51  0.6991492  43108.72
## 0.8         67.1298  36114.78  0.8006268  21929.86
## 0.8        155.0785  36082.93  0.8009233  21893.55
## 0.8        358.2515  36037.20  0.8013466  21822.43
## 0.8        827.6075  35965.21  0.8020353  21743.13
## 0.8       1911.8807  36290.07  0.7992346  21884.28
## 0.8       4416.6923  37253.84  0.7916940  22466.77
## 0.8      10203.1320  39588.76  0.7777147  24083.00
## 0.8      23570.5583  46287.64  0.7500972  29418.33
## 0.8      54451.0468  65720.45  0.6715741  46098.13
## 0.9         67.1298  36115.40  0.8006190  21930.51
## 0.9        155.0785  36077.55  0.8009745  21885.56
## 0.9        358.2515  36031.74  0.8013915  21813.24
## 0.9        827.6075  35969.58  0.8019961  21742.46
## 0.9       1911.8807  36419.57  0.7979926  21965.17
## 0.9       4416.6923  37442.26  0.7902141  22597.98
## 0.9      10203.1320  40114.96  0.7735845  24517.84
## 0.9      23570.5583  47699.38  0.7409547  30694.05
## 0.9      54451.0468  69218.41  0.6334724  49019.91
## 1.0         67.1298  36112.90  0.8006400  21927.45
## 1.0        155.0785  36072.70  0.8010177  21878.49
## 1.0        358.2515  36025.62  0.8014413  21806.18
## 1.0        827.6075  35985.43  0.8018439  21749.46
## 1.0       1911.8807  36560.02  0.7966230  22065.35
## 1.0       4416.6923  37643.93  0.7886044  22749.16
## 1.0      10203.1320  40652.79  0.7691381  24971.72
## 1.0      23570.5583  49252.63  0.7272645  32088.07
## 1.0      54451.0468  72575.49  0.6327540  51807.66
##
```

```
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were alpha = 0.1 and lambda =
4416.692.
```

```
modelNet$bestTune
```

```
##   alpha  lambda
## 6    0.1 4416.692
```

```
coef(modelNet$finalModel, modelNet$bestTune$lambda)
```

```
## 34 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              s1
## (Intercept) -446044.6606322
## MSSubClass  -124.1910671
## LotArea      0.3069953
## OverallQual 17454.8949706
## OverallCond 2975.4227571
## YearBuilt   259.5813826
## YearRemodAdd 210.0733854
## BsmtFinSF1  12.3885024
## BsmtFinSF2  -1.3857927
## BsmtUnfSF    .
## TotalBsmtSF 12.5431143
## X1stFlrSF    16.5328450
## X2ndFlrSF    13.8555191
## LowQualFinSF -27.6945759
## GrLivArea    27.8625508
## BsmtFullBath 6980.1065437
## BsmtHalfBath 219.3510046
## FullBath     4777.7752290
## HalfBath     .
## BedroomAbvGr -6868.6949722
## KitchenAbvGr -14607.7915368
## TotRmsAbvGrd 4558.5791748
## Fireplaces   4817.1568442
## GarageCars    9853.6589058
## GarageArea    8.1078208
## WoodDeckSF    27.5033252
## OpenPorchSF   .
## EnclosedPorch .
## X3SsnPorch    6.6175783
## ScreenPorch   34.7193789
## PoolArea     -19.8937409
## MiscVal       .
## MoSold        67.9057240
## YrSold        -268.3969512
```

Single Layer Neural Network

```
library(keras)
```

```
## Warning: package 'keras' was built under R version 4.2.3
```

```
library(tensorflow)
```

```
## Warning: package 'tensorflow' was built under R version 4.2.3
```

```
##
```

```
## Attaching package: 'tensorflow'
```

```

## The following object is masked from 'package:caret':
##
##      train

NNmodel <- keras_model_sequential () %>%
  layer_dense(units=50,activation="relu",input_shape=ncol(x)) %>%
  layer_dropout(rate=0.2) %>%
  layer_dense(units=1)

NNmodel %>% compile(loss="mse", optimizer = optimizer_adam(), metrics =
list("mean_absolute_error"))

#Fitting out NN model
history <- NNmodel %>% fit(x, y, epochs=50, batch_size=4,
validation_data=list(x_val,y_val))

## Epoch 1/50
## 322/322 - 1s - loss: 28281761792.0000 - mean_absolute_error: 149165.3281 -
val_loss: 13853976576.0000 - val_mean_absolute_error: 97784.9453 - 1s/epoch -
4ms/step
## Epoch 2/50
## 322/322 - 1s - loss: 9932403712.0000 - mean_absolute_error: 69713.3125 -
val_loss: 4693713920.0000 - val_mean_absolute_error: 46865.5625 - 713ms/epoch
- 2ms/step
## Epoch 3/50
## 322/322 - 1s - loss: 7184622080.0000 - mean_absolute_error: 50536.0117 -
val_loss: 3626127104.0000 - val_mean_absolute_error: 41690.0039 - 706ms/epoch
- 2ms/step
## Epoch 4/50
## 322/322 - 1s - loss: 6019992576.0000 - mean_absolute_error: 46959.5547 -
val_loss: 3340926720.0000 - val_mean_absolute_error: 40630.5078 - 722ms/epoch
- 2ms/step
## Epoch 5/50
## 322/322 - 1s - loss: 5437470720.0000 - mean_absolute_error: 46798.5625 -
val_loss: 3128451584.0000 - val_mean_absolute_error: 39899.2539 - 612ms/epoch
- 2ms/step
## Epoch 6/50
## 322/322 - 1s - loss: 4894033920.0000 - mean_absolute_error: 45847.8516 -
val_loss: 2949830656.0000 - val_mean_absolute_error: 39406.3086 - 583ms/epoch
- 2ms/step
## Epoch 7/50
## 322/322 - 1s - loss: 4596299776.0000 - mean_absolute_error: 44612.0820 -
val_loss: 2894654720.0000 - val_mean_absolute_error: 38460.5156 - 584ms/epoch
- 2ms/step
## Epoch 8/50
## 322/322 - 1s - loss: 4252326912.0000 - mean_absolute_error: 43478.0000 -
val_loss: 2752000768.0000 - val_mean_absolute_error: 38202.8125 - 610ms/epoch
- 2ms/step
## Epoch 9/50
## 322/322 - 1s - loss: 4059081728.0000 - mean_absolute_error: 43592.2656 -

```

```
val_loss: 2776863232.0000 - val_mean_absolute_error: 36964.5273 - 585ms/epoch  
- 2ms/step  
## Epoch 10/50  
## 322/322 - 1s - loss: 4024349184.0000 - mean_absolute_error: 42565.0156 -  
val_loss: 2595917824.0000 - val_mean_absolute_error: 37010.2578 - 660ms/epoch  
- 2ms/step  
## Epoch 11/50  
## 322/322 - 1s - loss: 3889306880.0000 - mean_absolute_error: 42232.5234 -  
val_loss: 2581872128.0000 - val_mean_absolute_error: 35645.6523 - 594ms/epoch  
- 2ms/step  
## Epoch 12/50  
## 322/322 - 1s - loss: 3951964928.0000 - mean_absolute_error: 42003.0195 -  
val_loss: 2536089344.0000 - val_mean_absolute_error: 35070.3906 - 591ms/epoch  
- 2ms/step  
## Epoch 13/50  
## 322/322 - 1s - loss: 3417050624.0000 - mean_absolute_error: 39489.0078 -  
val_loss: 2378597376.0000 - val_mean_absolute_error: 34986.9727 - 578ms/epoch  
- 2ms/step  
## Epoch 14/50  
## 322/322 - 1s - loss: 3436976128.0000 - mean_absolute_error: 40303.4336 -  
val_loss: 2353676032.0000 - val_mean_absolute_error: 33867.2969 - 611ms/epoch  
- 2ms/step  
## Epoch 15/50  
## 322/322 - 1s - loss: 3390011136.0000 - mean_absolute_error: 39942.0664 -  
val_loss: 2271399680.0000 - val_mean_absolute_error: 33037.0156 - 599ms/epoch  
- 2ms/step  
## Epoch 16/50  
## 322/322 - 1s - loss: 3297310720.0000 - mean_absolute_error: 38721.6875 -  
val_loss: 2178615040.0000 - val_mean_absolute_error: 32249.4512 - 591ms/epoch  
- 2ms/step  
## Epoch 17/50  
## 322/322 - 1s - loss: 3293226240.0000 - mean_absolute_error: 37852.7539 -  
val_loss: 2051465728.0000 - val_mean_absolute_error: 32370.3164 - 580ms/epoch  
- 2ms/step  
## Epoch 18/50  
## 322/322 - 1s - loss: 3059139840.0000 - mean_absolute_error: 38469.7305 -  
val_loss: 2085962880.0000 - val_mean_absolute_error: 30898.6191 - 609ms/epoch  
- 2ms/step  
## Epoch 19/50  
## 322/322 - 1s - loss: 3068960000.0000 - mean_absolute_error: 36358.8359 -  
val_loss: 1950469120.0000 - val_mean_absolute_error: 30504.9336 - 593ms/epoch  
- 2ms/step  
## Epoch 20/50  
## 322/322 - 1s - loss: 2978004480.0000 - mean_absolute_error: 37281.6367 -  
val_loss: 1836294912.0000 - val_mean_absolute_error: 30186.5469 - 590ms/epoch  
- 2ms/step  
## Epoch 21/50  
## 322/322 - 1s - loss: 3092871680.0000 - mean_absolute_error: 37347.8125 -  
val_loss: 1793910656.0000 - val_mean_absolute_error: 29326.5371 - 587ms/epoch  
- 2ms/step
```



```
## Epoch 22/50
## 322/322 - 1s - loss: 2836470016.0000 - mean_absolute_error: 35907.5977 -
val_loss: 1772911744.0000 - val_mean_absolute_error: 28845.6504 - 601ms/epoch
- 2ms/step
## Epoch 23/50
## 322/322 - 1s - loss: 2844355328.0000 - mean_absolute_error: 35422.7539 -
val_loss: 1720975872.0000 - val_mean_absolute_error: 28355.4590 - 584ms/epoch
- 2ms/step
## Epoch 24/50
## 322/322 - 1s - loss: 2775760384.0000 - mean_absolute_error: 35349.6289 -
val_loss: 1677808256.0000 - val_mean_absolute_error: 27806.7031 - 585ms/epoch
- 2ms/step
## Epoch 25/50
## 322/322 - 1s - loss: 2840767232.0000 - mean_absolute_error: 35431.0664 -
val_loss: 1637430912.0000 - val_mean_absolute_error: 27498.0781 - 587ms/epoch
- 2ms/step
## Epoch 26/50
## 322/322 - 1s - loss: 2734164736.0000 - mean_absolute_error: 35024.2305 -
val_loss: 1644825216.0000 - val_mean_absolute_error: 27321.3770 - 585ms/epoch
- 2ms/step
## Epoch 27/50
## 322/322 - 1s - loss: 2629021696.0000 - mean_absolute_error: 34225.1328 -
val_loss: 1583942912.0000 - val_mean_absolute_error: 27305.3770 - 580ms/epoch
- 2ms/step
## Epoch 28/50
## 322/322 - 1s - loss: 2889126912.0000 - mean_absolute_error: 35627.6641 -
val_loss: 1587999744.0000 - val_mean_absolute_error: 27230.9863 - 588ms/epoch
- 2ms/step
## Epoch 29/50
## 322/322 - 1s - loss: 2684827648.0000 - mean_absolute_error: 34526.8789 -
val_loss: 1520063872.0000 - val_mean_absolute_error: 27310.4062 - 604ms/epoch
- 2ms/step
## Epoch 30/50
## 322/322 - 1s - loss: 2755955712.0000 - mean_absolute_error: 34549.2695 -
val_loss: 1554321280.0000 - val_mean_absolute_error: 26938.0059 - 584ms/epoch
- 2ms/step
## Epoch 31/50
## 322/322 - 1s - loss: 2882132992.0000 - mean_absolute_error: 35165.8555 -
val_loss: 1504360448.0000 - val_mean_absolute_error: 26680.8984 - 594ms/epoch
- 2ms/step
## Epoch 32/50
## 322/322 - 1s - loss: 2802217984.0000 - mean_absolute_error: 34216.1719 -
val_loss: 1544844288.0000 - val_mean_absolute_error: 27031.9395 - 587ms/epoch
- 2ms/step
## Epoch 33/50
## 322/322 - 1s - loss: 2714020352.0000 - mean_absolute_error: 33967.0156 -
val_loss: 1513739520.0000 - val_mean_absolute_error: 26831.5020 - 598ms/epoch
- 2ms/step
## Epoch 34/50
## 322/322 - 1s - loss: 2704364288.0000 - mean_absolute_error: 34183.5508 -
```


```
val_loss: 1517229440.0000 - val_mean_absolute_error: 26922.4121 - 584ms/epoch  
- 2ms/step  
## Epoch 35/50  
## 322/322 - 1s - loss: 2674246400.0000 - mean_absolute_error: 34617.4258 -  
val_loss: 1474669952.0000 - val_mean_absolute_error: 26749.0840 - 582ms/epoch  
- 2ms/step  
## Epoch 36/50  
## 322/322 - 1s - loss: 2755875584.0000 - mean_absolute_error: 34366.6367 -  
val_loss: 1469974016.0000 - val_mean_absolute_error: 26770.6367 - 645ms/epoch  
- 2ms/step  
## Epoch 37/50  
## 322/322 - 1s - loss: 2588229376.0000 - mean_absolute_error: 33527.0625 -  
val_loss: 1576558464.0000 - val_mean_absolute_error: 27586.8730 - 584ms/epoch  
- 2ms/step  
## Epoch 38/50  
## 322/322 - 1s - loss: 2660746496.0000 - mean_absolute_error: 33934.0820 -  
val_loss: 1527195136.0000 - val_mean_absolute_error: 27324.3867 - 581ms/epoch  
- 2ms/step  
## Epoch 39/50  
## 322/322 - 1s - loss: 2812005376.0000 - mean_absolute_error: 34100.0312 -  
val_loss: 1451259520.0000 - val_mean_absolute_error: 26785.5078 - 584ms/epoch  
- 2ms/step  
## Epoch 40/50  
## 322/322 - 1s - loss: 2734551808.0000 - mean_absolute_error: 34008.7188 -  
val_loss: 1505810304.0000 - val_mean_absolute_error: 27251.5254 - 590ms/epoch  
- 2ms/step  
## Epoch 41/50  
## 322/322 - 1s - loss: 2766029568.0000 - mean_absolute_error: 34267.1797 -  
val_loss: 1515094912.0000 - val_mean_absolute_error: 27357.5527 - 584ms/epoch  
- 2ms/step  
## Epoch 42/50  
## 322/322 - 1s - loss: 2648256512.0000 - mean_absolute_error: 33872.7969 -  
val_loss: 1477571968.0000 - val_mean_absolute_error: 27098.9980 - 587ms/epoch  
- 2ms/step  
## Epoch 43/50  
## 322/322 - 1s - loss: 2590374400.0000 - mean_absolute_error: 34580.2227 -  
val_loss: 1425483776.0000 - val_mean_absolute_error: 26700.3633 - 610ms/epoch  
- 2ms/step  
## Epoch 44/50  
## 322/322 - 1s - loss: 2564006144.0000 - mean_absolute_error: 33786.2188 -  
val_loss: 1434891392.0000 - val_mean_absolute_error: 26883.0059 - 585ms/epoch  
- 2ms/step  
## Epoch 45/50  
## 322/322 - 1s - loss: 2592933376.0000 - mean_absolute_error: 34074.8281 -  
val_loss: 1457150464.0000 - val_mean_absolute_error: 27077.7676 - 604ms/epoch  
- 2ms/step  
## Epoch 46/50  
## 322/322 - 1s - loss: 2945060352.0000 - mean_absolute_error: 34758.0156 -  
val_loss: 1463529984.0000 - val_mean_absolute_error: 27087.9121 - 687ms/epoch  
- 2ms/step
```

```

## Epoch 47/50
## 322/322 - 1s - loss: 2694374144.0000 - mean_absolute_error: 33500.2734 -
val_loss: 1503827840.0000 - val_mean_absolute_error: 27375.2129 - 623ms/epoch
- 2ms/step
## Epoch 48/50
## 322/322 - 1s - loss: 2683096320.0000 - mean_absolute_error: 34446.2930 -
val_loss: 1435055744.0000 - val_mean_absolute_error: 26869.6484 - 579ms/epoch
- 2ms/step
## Epoch 49/50
## 322/322 - 1s - loss: 2891406080.0000 - mean_absolute_error: 35096.5781 -
val_loss: 1510734464.0000 - val_mean_absolute_error: 27511.3809 - 581ms/epoch
- 2ms/step
## Epoch 50/50
## 322/322 - 1s - loss: 2672335104.0000 - mean_absolute_error: 33863.3828 -
val_loss: 1529366912.0000 - val_mean_absolute_error: 27670.3730 - 597ms/epoch
- 2ms/step

```

`plot(history)`



Model Comparisons

So we have created 5 models: A regular linear regression with all our columns, another regular linear regression with strongly correlated columns, a Lasso regression model, An ElasticNet regression model, and a simple Neural Network model.

One thing we notice in our neural network model is that the validation loss and validation absolute error is less than the training set equivalent. This could have been due to our dropout rate which reduced the overfitting in our training set.

Let's compare R squared values in our regression models.

Regular regression R²

```
summary(model_orig)
```

```
##
## Call:
## lm(formula = SalePrice ~ MSSubClass + LotArea + OverallQual +
##      OverallCond + YearBuilt + BsmtFinSF1 + YearRemodAdd + BsmtFinSF2 +
##      BsmtUnfSF + TotalBsmtSF + X1stFlrSF + X2ndFlrSF + LowQualFinSF +
##      GrLivArea + BsmtFullBath + BsmtHalfBath + FullBath + HalfBath +
##      BedroomAbvGr + KitchenAbvGr + TotRmsAbvGrd + Fireplaces +
##      GarageCars + GarageArea + WoodDeckSF + OpenPorchSF + EnclosedPorch +
##      X3SsnPorch + ScreenPorch + PoolArea + MiscVal + MoSold +
##      YrSold, data = train.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -475511  -16893   -2383   13380   289170
##
## Coefficients: (2 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  345266.9063 1557437.0846   0.222   0.824592
## MSSubClass    -155.9361    28.9110  -5.394 0.0000000824221369 ***
## LotArea         0.3463     0.1060   3.266   0.001119 **
## OverallQual   18308.4077   1302.3680  14.058 < 0.00000000000000002 ***
## OverallCond    4094.3065   1109.2275   3.691   0.000233 ***
## YearBuilt      326.0210     66.7722   4.883 0.0000011816567526 ***
## BsmtFinSF1     22.8103     5.0835   4.487 0.0000078803809301 ***
## YearRemodAdd   149.5174     71.4315   2.093   0.036535 *
## BsmtFinSF2      5.9618     7.8855   0.756   0.449763
## BsmtUnfSF     11.6895     4.5968   2.543   0.011111 *
## TotalBsmtSF         NA          NA      NA      NA
## X1stFlrSF      48.4809     6.2356   7.775 0.00000000000000156 ***
## X2ndFlrSF      48.2298     5.4493   8.851 < 0.00000000000000002 ***
## LowQualFinSF   -6.1005    24.5491  -0.249   0.803788
## GrLivArea         NA          NA      NA      NA
## BsmtFullBath    9255.4046   2862.1693   3.234   0.001254 **
## BsmtHalfBath    3378.1330   4442.1203   0.760   0.447112
## FullBath       3830.4355   3105.9303   1.233   0.217708
## HalfBath       -1375.9813   2926.9289  -0.470   0.638357
## BedroomAbvGr  -10194.7654   1857.2890  -5.489 0.0000000488419628 ***
## KitchenAbvGr  -14752.3253   5692.7753  -2.591   0.009669 **
## TotRmsAbvGrd   5378.7065   1350.4815   3.983 0.0000720153832422 ***
## Fireplaces     3587.7710   1921.2523   1.867   0.062077 .
## GarageCars     10962.9651   3120.5299   3.513   0.000459 ***
```

```
## GarageArea      -0.8172      10.6759  -0.077      0.938998
## WoodDeckSF      28.5755       8.6994   3.285      0.001049 **
## OpenPorchSF     -14.3054      16.2529  -0.880      0.378932
## EnclosedPorch    9.8982      18.3120   0.541      0.588926
## X3SsnPorch      17.7876      33.2294   0.535      0.592539
## ScreenPorch     47.8275      18.8194   2.541      0.011160 *
## PoolArea       -36.0041      24.4594  -1.472      0.141272
## MiscVal         -0.4850       1.9282  -0.252      0.801429
## MoSold          186.9684     376.4723   0.497      0.619535
## YrSold          -671.7207     774.8394  -0.867      0.386154
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 35920 on 1256 degrees of freedom
## Multiple R-squared:  0.8013, Adjusted R-squared:  0.7964
## F-statistic: 163.4 on 31 and 1256 DF,  p-value: < 0.00000000000000022
```

`summary(model_reduced)`

```
##
## Call:
## lm(formula = SalePrice ~ OverallQual + YearBuilt + BsmtFinSF1 +
##      YearRemodAdd + TotalBsmtSF + X1stFlrSF + X2ndFlrSF + GrLivArea +
##      FullBath + TotRmsAbvGrd + Fireplaces + GarageCars + GarageArea +
##      WoodDeckSF + OpenPorchSF, data = train.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -511461  -17382   -2175   14405   290530
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1081490.742  137374.324  -7.873  0.00000000000000739 ***
## OverallQual    20019.147   1271.679  15.742 < 0.0000000000000002 ***
## YearBuilt      175.488     54.301    3.232   0.001262 **
## BsmtFinSF1     19.117      2.819    6.781  0.000000000001821744 ***
## YearRemodAdd   333.681     66.845    4.992  0.00000068114781674 ***
## TotalBsmtSF    13.441      4.665    2.881   0.004029 **
## X1stFlrSF      57.612     25.526    2.257   0.024177 *
## X2ndFlrSF      50.108     25.173    1.991   0.046748 *
## GrLivArea     -11.206     25.050   -0.447   0.654709
## FullBath     -1674.446   2848.891   -0.588   0.556802
## TotRmsAbvGrd   2081.566   1176.934    1.769   0.077195 .
## Fireplaces     7268.603   1918.762    3.788   0.000159 ***
## GarageCars    10164.873   3209.093    3.168   0.001574 **
## GarageArea     10.807     10.966    0.986   0.324541
## WoodDeckSF      35.645      8.800    4.051  0.00005414055959256 ***
## OpenPorchSF    -3.722     16.658   -0.223   0.823216
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 37340 on 1272 degrees of freedom
## Multiple R-squared:  0.7826, Adjusted R-squared:  0.78
## F-statistic: 305.2 on 15 and 1272 DF,  p-value: < 0.00000000000000022
```

Surprisingly, our linear model that has every column has a better R^2 than the select columns with high correlation. This may be due to too many features lost which underfit the training data.

Lasso and ElasticNet R^2 and RMSE

Firstly using lasso and ElasticNet, we make predictions on our validation set and compare the RMSE as well

```
# LASSO
predictionLasso <- modelLass %>% predict(x_val)

# RMSE and R^2
data.frame(
  RMSE = caret::RMSE(predictionLasso, y_val),
  Rsquare = caret::R2(predictionLasso, y_val))

##          RMSE          s0
## 1 30717.86 0.8497609

# ELASTIC NET
predictionNet <- modelNet %>% predict(x_val)

# RMSE and R^2
data.frame(
  RMSE = caret::RMSE(predictionNet, y_val),
  Rsquare = caret::R2(predictionNet, y_val))

##          RMSE    Rsquare
## 1 29225.56 0.8625493
```

As we can see, out of all of our regression models, ElasticNet has the largest R squared value meaning it better fits our training data. We also notice the RMSE of ElasticNet is lower than Lasso which concludes that it is more accurate.

When viewing the coefficients of ElasticNet above, we see that the intercept $\hat{\beta}_0$ is -446044. This model also gives us predictors for multiple other columns but also does not include some such as BsmtUnfSF.

Therefore out of our regression models, ElasticNet seems to be the best performing.

ElasticNet vs Neural Network

We will now see the final predictions of our ElasticNet regression model, and compare to the neural network model we made. We will submit to kaggle and see our score on the testing data.

```
elastic <- modelNet %>% predict(test.data)
elastic <- cbind(Id = test_dataa$Id, elastic)
colnames(elastic)[colnames(elastic) == 'elastic'] <- 'SalePrice'

NN <- predict(NNmodel, test.data)

## 46/46 - 0s - 125ms/epoch - 3ms/step

NN <- cbind(Id = test_dataa$Id, NN)
colnames(NN)[colnames(NN) == ''] <- 'SalePrice'

write.csv(elastic, "elastic.csv", row.names = FALSE)

write.csv(NN, "nn.csv", row.names = FALSE)
```

Results

After Submitting, it turns out our Neural Network model is the best performing model, beating out all the linear regression variants. The score for the NN model on Kaggle was 0.22206 while the ElasticNet score was 0.34880. Both are very solid scores but the NN model is the best. This may be due to the number of epochs that I have ran for the NN model, 50 is a lot.

Therefore, overall the best regression model was our ElasticNet model, but it has been beaten by a single layer neural network.

For future model improvements, some heavy feature selection can be made while also testing different parameters like the alpha values in regression.